

Homework 1: Programming Component

Using Minimum Spanning Trees to Route Packets in a Network Overlay

VERSION 1.0

DUE DATE: Wednesday September 24th, 2025 @ 8:00 pm

The objective of this assignment is to get you familiar with coding in a distributed setting where you need to manage the underlying communications between nodes. Upon completion of this assignment you will have a set of reusable classes that you will be able to draw upon. You are required to work alone on this assignment. Use of GenAI tools is expressly prohibited; see the textbox below. The assignment accounts for **10% towards your cumulative course grade**. This assignment may be modified to clarify any questions (and the version number incremented), but the crux of the assignment and the distribution of points will not change. If there any changes to the assignment, all changes will be document in the "Change History" section of this assignment.

Generative AI Use and Consequences

Use of AI tools such as ChatGPT, Claude, Github Co-Pilot, or anything of their kind to write or "improve" your code or written work at *any* stage is prohibited; this includes the ideation phase. It is your responsibility to ensure that you don't have the GitHub Co-Pilot extension installed in your IDE; assignment solutions generated by Co-Pilot aren't written by you. Turning in code or an essay written by generative AI tools will be treated as turning in work created by someone else, namely an act of plagiarism and/or cheating. At a minimum, this will result in a 100% deduction (i.e., you will receive -10/10). To ensure fairness and maintain integrity, grading will also include code reviews, interviews, and on-the-spot code modifications.

Ultimately, you will get out of the class what you put in. Simply copying and pasting code from generative AI tools is not only unethical, it robs you of the chance to learn. Here are four reasons why these generative AI tools undercuts your own education:

1. They take away the struggle that leads to understanding. They rob you of the ability to think and learn the concepts for yourself. Solving problems yourself is how concepts stick. If the AI does the work, what's left for you to learn?
2. You will struggle with the in-classroom quizzes and exams where you will not have access to these tools.
3. Yes, AI tools will become an important part of a software engineer's workflow. But to use them effectively later, you first need solid expertise in the subject matter; and, that only comes from practicing *without* them.
4. These tools are prone to generating imperfect or even incorrect solutions, so trusting them blindly can lead to bad consequences.

Crux of the Assignment

As part of this assignment you will be: (1) constructing a logical overlay over a distributed set of nodes, and then (2) constructing minimum spanning trees to route packets in the system.

The overlay will contain at least 10 messaging nodes, and each messaging node will be connected to **C_R** (default of 4) other messaging nodes. Each link that connects two messaging nodes within the overlay has a weight associated with it. Links are **bidirectional** i.e. if messaging node **A** established a connection to messaging node **B**, then messaging node **B** must use that link to communicate with **A**.

Once the overlay has been set up, messaging nodes in the system will select a node at random and send that node (also known as the sink node) a message. Rather than send this message directly to the sink node, the source node will use the overlay for communications. This is done by computing a Minimum Spanning Tree (MST), rooted at the source node, based on the link weights assigned during overlay construction. The MST is a subset of the network's edges that connects all nodes together without cycles and with the smallest possible total link weight.

Because the MST is built from the perspective of the source, each source node will typically produce a different routing tree. The path from the source to the sink is determined by following the edges of this MST, which may involve zero or more intermediate messaging nodes that relay the packets. Such intermediate nodes are said to **relay** the packets.

The assignment requires you to verify the correctness of packet exchanges between sources and sinks by ensuring that: (1) the number of messages sent and received within the system match, and (2) these messages have not been corrupted in transit to the intended recipient.

All communications in this assignment are based on **TCP**. The assignment must be implemented in **Java** and you cannot use any external jar files. You must develop all functionality yourself.

1 Components

There are two components that you will be building as part of this assignment: a registry and a messaging node.

1.1 Registry:

There is exactly one registry in the system. The registry provides the following functions:

- A. Allows messaging nodes to register themselves. This is performed when a messaging node starts up for the first time.
- B. Allows messaging nodes to deregister themselves. This is performed when a messaging node leaves the overlay.
- C. Enables the construction of the overlay by orchestrating connections that a messaging node initiates with other messaging nodes in the system. Based on its knowledge of the messaging nodes (through function A) the registry informs messaging nodes about the other messaging nodes that they should connect to.
- D. Assign and publish weights to the links connecting any two messaging nodes in the overlay. The weights these links take will range from 1-10.

The registry maintains information about the registered messaging nodes in a registry; you can use any data structure for managing this registry but make sure that your choice can support all the operations that you will need.

The registry does not play any role in the *routing* of data within the overlay. Interactions between the messaging nodes and the registry are via request-response messages. For each request that it receives from the messaging nodes, the registry will send a response back to the messaging node (based on the IP address associated with Socket's input stream) where the request originated. The contents of this response depend on the *type* of the request and the *outcome* of processing this request.

1.2 The Messaging node

Unlike the registry, there are multiple messaging nodes (minimum of 10) in the system. A messaging node provides two closely related functions: it initiates and accepts both communications and messages within the system.

Communications that nodes have with each other are based on TCP. Each messaging node needs to automatically configure the ports over which it listens for communications i.e. the port numbers should not be hard-coded or specified at the command line. `TCPServerSocket` is used to accept incoming TCP communications.

Once the initialization is complete, the node should send a registration request to the registry.

2 Interactions between the components

This section will describe the interactions between the registry and the messaging nodes. This section includes the prescribed wire-formats. You have freedom to construct your wire-formats but you must include the fields that have been specified. A good programming practice is to have a separate class for each message type so that you can isolate faults better. The `Message Types` that have been specified could be part of an interface, say `csx55.overlay.wireformats.Protocol` and have values specified there. This way you are not hard-coding values in different portions of your code.

Use of Java serialization is not allowed. Your classes for the message types should not implement the `java.io.Serializable` interface.

2.1 Registration:

Upon starting up, each messaging node should register its IP address, and port number with the registry. It should be possible for your system to register messaging nodes that are running on the same host but are listening to communications on different ports. There should be 3 fields in this registration request:

```
Message Type (int): REGISTER_REQUEST
IP address (String)
Port number (int)
```

When a registry receives this request, it checks to see if the node had previously registered and ensures the IP address in the message matches the address where the request originated. The registry issues an error message under two circumstances:

- If the node had previously registered and has a valid entry in its registry.
- If there is a mismatch in the address that is specified in the registration request and the IP address of the request (the socket's input stream).

The contents of the response message are depicted below. The success or failure of the registration request should be indicated in the status field of the response message.

```
Message Type (int): REGISTER_RESPONSE
Status Code (byte): SUCCESS or FAILURE
Additional Info (String):
```

In the case of successful registration, the registry should include a message that indicates the number of entries currently present in its registry. A sample information string is "Registration request successful. The number of messaging nodes currently constituting the overlay is (5)". If the registration was unsuccessful, the message from the registry should indicate why the request was unsuccessful.

NOTE: In the rare case that a messaging node fails just after it sends a registration request, the registry will not be able to communicate with it. In this case, the entry for the messaging node should be removed from the messaging node-registry maintained at the registry.

2.2 Deregistration

When a messaging node exits it should deregister itself. It does so by sending a message to the registry. This deregistration request includes the following fields

```
Message Type: Deregister_Request
Node IP address:
Node Port number:
```

The registry should check to see that request is a valid one by checking (1) where the message originated and (2) whether this node was previously registered. Error messages should be returned in case of a mismatch in the addresses or if the messaging node is not registered with the overlay. You should be able to test the error-reporting functionality by de-registering the same messaging node twice.

2.3 Peer messaging nodes list

Once the **setup-overlay** command (see section 3) is specified at the registry it must perform a series of actions that lead to the creation of the overlay via messaging nodes initiating connections with each other. Messaging nodes await instructions from the registry regarding the other messaging nodes that they must establish connections to.

The registry must ensure two properties. First, it must ensure that the number of links to/from (the links are bidirectional) every messaging node in the overlay is identical; this is configurable metric (with a default value of 4) and is specified as part of the **setup-overlay** command. Second, the registry must ensure that there is *no partition* within the overlay i.e. it should be possible to reach any messaging node from any other messaging node in the overlay.

If the connection requirement for the overlay is **C_R**, each messaging node will have **C_R** links to other messaging nodes in the overlay. The registry selects these **C_R** messaging nodes that constitute the *peer-messaging nodes list* for a messaging node randomly. However, a check should be performed to ensure that the peer-messaging nodes list for a messaging node does not include the targeted messaging node i.e. a messaging node should not have to connect to itself. The registry keeps track of the connections that are being created; for example, if messaging node A is asked to connect to messaging node B, the connection counts for both A and B are incremented. The registry must ensure that connection counts are met and not breached.

The registry sends a different list of messaging nodes to each messaging node in the overlay. To avoid duplicate connections being established between messaging nodes, only one messaging node in a link should be instructed to create the connection. For instance, if there is a link between nodes A and B, only node A should be instructed to establish a link with node B or vice versa. The number of peer messaging nodes included in messages to different messaging nodes may vary from **C_R** through 0. If a messaging node's connection limit was reached through previous messages sent to other messaging nodes in the overlay, a message still needs to be sent to that messaging node. The peer-list message will have the following format

```
Message Type: Messaging_Nodes_List
Number of peer messaging nodes: X
Messaging node1 Info
Messaging node2 Info
...
Messaging nodeX Info
```

If a messaging node does not need to establish any connections, set the number of peer messaging nodes to 0. The information corresponding to a messaging node includes the following: `messaging node_ip:portnum`. Upon receiving the `MESSAGING_NODES_LIST` message a messaging node should initiate connections to the specified messaging nodes. After establishing connections, a messaging node should print the message "All connections are established. Number of connections: x" to the console to help with testing and evaluating your implementation.

2.4 Assign overlay link weights

The registry is also responsible for assigning weights to connections in the overlay. The weight for each link is an integer between 1-10 and is randomly computed by the registry. This information will be encoded in the message as follows.

```
Message Type: Link_Weights
Number of links: L
Linkinfo1
Linkinfo2
...
LinkinfoL
```

A `Linkinfo` connecting messaging nodes A and B contains the following fields: `ipA:portnumA`
`ipB:portnumB` `weight`

A single message should be constructed with all link weights and sent to all registered messaging nodes. A messaging node should process this message and store its information to generate routing paths for messages as explained in the following section. Further, it should acknowledge the receipt and processing of this message by printing the message "Link weights received and processed. Ready to send messages." to the console.

2.5 Initiate sending messages

The registry informs nodes in the overlay when they should start sending messages to each other. It does so via the `TASK_INITIATE` control message.

```
Message Type: TASK_INITIATE
Rounds: X
```

2.6 Send message

Data can be fed into the network from any messaging node within the network. Packets are sent from a source to a sink; it is possible that there might be zero or more intermediate nodes in the system that **relay** messages *en route* to the sink. Every node tracks the number of messages that it has relayed during communications within the overlay.

When a packet is ready to be sent from a source node to the sink node, the source node first constructs a Minimum Spanning Tree (MST) for the network using an algorithm such as Prim's or Kruskal's. An MST is a subset of the network's edges that connects all the nodes together without cycles and with the smallest possible total link weight. It provides a unique, loop-free structure for routing messages.

In this assignment, the MST is built from the perspective of the source node—that is, the source treats itself as the root and computes a tree that spans all other nodes in the network. Because each node builds its own MST when acting as the source, the structure and routing paths will generally differ from node to node. The resulting MST is then used as the routing plan: each packet is forwarded along the edges of that tree toward the sink node. For example, although A may have a direct connection to B, the MST rooted at A may instead route the packet along $A \rightarrow C \rightarrow E \rightarrow D \rightarrow B$ if that path is part of the spanning tree. Since the link weights are fixed once assigned, you may cache the MST once computed as an optimization. Also, a message includes a payload, which is a random integer as explained in Section 4.

A key requirement for the dissemination of packets within the overlay is that no messaging node should receive the same packet more than once. This should be achieved without having to rely on duplicate detection and suppression.

2.7 Inform registry of task completion

Once a node has completed its task of sending a certain number of messages in *rounds* (described in [section 4](#)), it informs the registry of its task completion using the `TASK_COMPLETE` message. This message will have the following format:

```
Message Type: TASK_COMPLETE
Node IP address:
Node Port number:
```

2.8 Retrieve traffic summaries from nodes

Once the registry has received `TASK_COMPLETE` messages from all the registered nodes it will issue a `PULL_TRAFFIC_SUMMARY` message. This message is sent to all the registered nodes in the system. This message will have the following format. To allow all messages that are already in transit to reach their destination nodes, you should wait for some time (e.g., 15 seconds) after receiving all `TASK_COMPLETE` messages before issuing a `PULL_TRAFFIC_SUMMARY` message.

```
Message Type: PULL_TRAFFIC_SUMMARY
```

2.9 Sending traffic summaries from the nodes to the registry

Upon receipt of the `PULL_TRAFFIC_SUMMARY` message from the registry, the node will create a response that includes summaries of the traffic that it has participated in. The summary will include information about messages that were sent and received. This message will have the following format.

```
Message Type: TRAFFIC_SUMMARY
Node IP address:
Node Port number:
Number of messages sent
Summation of sent messages
Number of messages received
Summation of received messages
Number of messages relayed
```

Once the `TRAFFIC_SUMMARY` message is sent to the registry, the node must reset the counters associated with traffic relating to the messages it has sent and received so far e.g number of messages sent, summation of sent messages, etc.

3 Specifying commands and interacting with the processes

Both the registry and the messaging node should run as *foreground* processes and allow support for commands to be specified while the processes are running. The commands that should be supported are specific to the two components. Always follow the format given for the output of the commands.

3.1 Registry

list-messaging-nodes

This should result in information about the messaging nodes (IP, and port-number) being listed. Information for each messaging node should be listed on a separate line; see an example below:

```
192.168.0.10:8080
192.168.1.25:443
192.168.100.5:22
10.0.0.15:3306
10.1.2.3:5000
10.10.10.50:21
172.16.0.20:8081
172.16.5.100:1433
172.20.10.2:25
172.31.255.200:27017
```

list-weights

This should list information about links comprising the overlay. Each link's information should be on a separate line and include information about the nodes that it connects to and the weight of that link. For example:

```
192.168.0.10:8080, 192.168.1.25:443, 8
```

indicates that the link is between two messaging nodes (192.168.0.10:8080 and 192.168.1.25:443) with a link weight of 8.

Print the full list:

```
192.168.0.10:8080, 192.168.1.25:443, 8
192.168.0.10:8080, 192.168.100.5:22, 5
172.16.5.100:1433, 192.168.1.25:443, 13
```

setup-overlay <number-of-connections>

This should result in the registry setting up the overlay. It does so by sending messaging nodes messages containing information about the messaging nodes that it should connect to. The registry tracks the connection counts for each messaging node and will send the `MESSAGING_NODES_LIST` message (see Section 2.3) to every messaging node. A sample specification of this command is **setup-overlay 4** that will result in the creation of an overlay where each messaging node is connected to exactly 4 other messaging nodes in the overlay. You should handle the error condition where the number of messaging nodes is less than the connection limit that is specified.

NOTE: You are not required to deal with the case where a messaging node is added or removed after the overlay has been set up. You must however deal with the case where a messaging node registers and deregisters from the registry before the overlay is set up.

The output of this command should just be the string "setup completed with <n> connections", with the correct number of connections, followed by the table with the results as explained in section 4.3.

send-overlay-link-weights

This should result in a `Link_Weights` message being sent to all registered nodes in the overlay. This command is issued once after the **setup-overlay** command has been issued. This also allows all nodes in the system to be aware of not just all the nodes in the system, but also the complete set of links in the system.

The following figure depicts an example overlay with 8 messaging nodes with the connection requirement (C_R) set to 4. For instance, node A is connected with nodes B, C, G and H using bidirectional links with weights of 1, 4, 7 and 6 respectively.




Figure 1: Graphical depiction of an overlay and link weights.

The output of this command should just be the string "link weights assigned"

start <number-of-rounds>

The **start** command results in nodes exchanging messages within the overlay. Each node in the overlay will be responsible for sending **number-of-rounds** messages. An advantage of this is that you are able to debug your system with a smaller set of messages and verify correctness of your programs across a wide range of values. A detailed description is provided in **section 4** below.

Once all rounds have been completed, the registry should print "<n> rounds completed", with the correct number of rounds

3.2 Messaging node

print-mst

This should print the MST that has been computed from the perspective of the current messaging node (i.e., with this node as the root). The MST should be listed in Breadth-First Search (BFS) order starting from the root. For each edge, display the two connected nodes and the weight of the link between them. This is the same format as the command "list-weights". For example:

Print the full list:

```
192.168.0.10:8080, 192.168.1.25:443, 8
192.168.0.10:8080, 192.168.100.5:22, 5
172.16.5.100:1433, 192.168.1.25:443, 13
```

exit-overlay

This allows a messaging node to exit the overlay. The messaging node should first send a deregistration message (see Section 2.2) to the registry and await a response before exiting and terminating the process. The output of this command should be "exited overlay".

4 Setting

For the remainder of the discussion, we assume that the `setup-overlay` command has been specified followed by the `send-overlay-link-weights` command at the registry. Also, nodes will not be added to the system from hereon.

When the `start` command is specified at the registry, the registry sends the `TASK_INITIATE` control message to all the registered nodes within the overlay. Upon receiving this information from the registry, a given node will start exchanging messages with other nodes.

Each node participates in a set of *rounds*. Each round involves a node sending 5 messages to a randomly chosen node (excluding itself, of course) from the set of registered nodes based on the `Link_Weights` message. All communications in the system will be based on TCP. To send a message the source node computes a routing plan (encoding the shortest path) with zero or more intermediate nodes relaying the message en route to the destination sink node. The payload of each message is a random integer with values that range from 2147483647 to -2147483648. At the end of each round, the process is repeated by choosing another node at random. The number of rounds initiated by each node is determined by the specified `number-of-rounds`.

The number of nodes will be fixed at the start of the experiment. We will likely use around 10 nodes for the test environment during grading. When setting up the overlay, a messaging node only opens one connection at a time to another node, it may receive multiple incoming connections as other nodes try to connect to it.

4.1 Tracking communications between nodes

Each node will maintain two integer variables that are initialized to zero: `sendTracker` and `receiveTracker`. The `sendTracker` represents the number of messages that were sent by that node and the `receiveTracker` maintains information about the number of messages that were received. Additionally, each node will track the number of messages that it relayed – i.e., messages for which it was neither the source nor the sink. Consider the case where there are 10 nodes in the system as depicted in Figure 1. Since every node initiates 5000 rounds, each of which contains 5 messages, the number of messages sent by every node is 25,000. With 10 nodes in the system, the total number of messages would be 250,000. Since a sending node chooses the target node for each round at random, the number of messages received by different receivers would be different. However, because each round has 5 messages, the total number of messages received at a receiver would be a multiple of 5 and close to 25,000 (i.e. it could be 24000, 24595, 26905, ...).

The number of messages that a node relays will depend on the overlay topology, link weights, and shortest paths that it belongs to. This is tracked using the variable `relayTracker`.

To track the messages that it has sent and received, each node will maintain two additional `long` variables that are initialized to zero: `sendSummation` and `receiveSummation`. The data type for these variables is a `long` to cope with overflow issues that will arise as part of the summing operations that will be performed. The variable `sendSummation`, continuously sums the values of the random numbers that are sent, while the `receiveSummation` sums values of the payloads that are received. The values of `sendSummation` and `receiveSummation` at a node can be positive or negative.

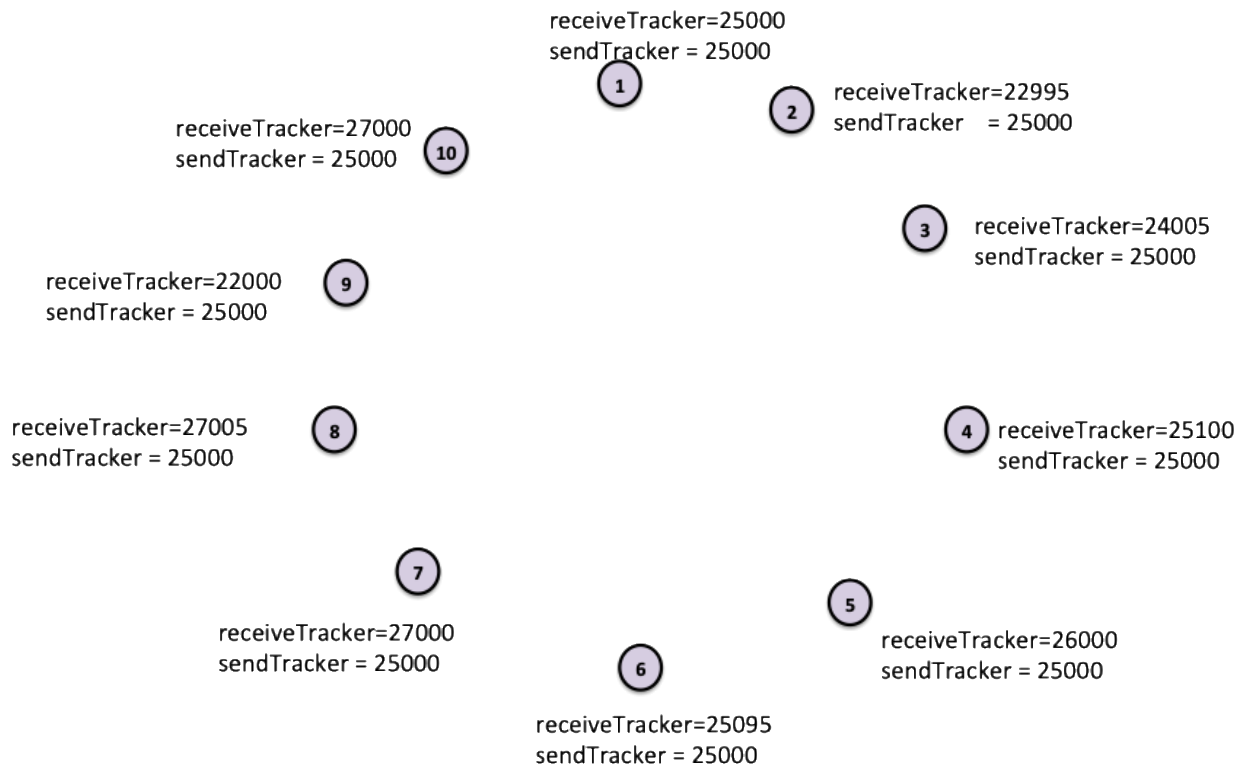


Figure 2: Depiction of a possible distribution of the number of messages sent and received within a set of 10 nodes.

4.2 Correctness Verification

We will verify **correctness** by: (1) checking the number of messages that were sent and received, and (2) if these packets were corrupted for some reason.

The total number of messages that were sent and received by the set of all nodes must match i.e. the cumulative sum of the `receiveTracker` at each node must match the cumulative sum of the `sendTracker` variable at each node. We will check that these packets were not corrupted by verifying that when we add up the values of `sendSummation` it will exactly match the added up values of `receiveSummation`.

4.3 Collecting and printing outputs

When a node has completed its rounds, it will send a `TASK_COMPLETE` message to the registry. When the registry receives a `TASK_COMPLETE` message from each of the N registered nodes in the system, it issues a `PULL_TRAFFIC_SUMMARY` message to all the nodes.

Upon receipt of the `PULL_TRAFFIC_SUMMARY` message, a node will prepare to send information about the messages that it has sent and received. This includes: (1) the number of messages that were sent by that node, (2) the summation of the sent messages, (3) the number of messages that were received by that node, (4) the summation of the received messages, and (5) the number of messages relayed by that node. The node packages this information in the `TRAFFIC_SUMMARY` message and sends it to the registry. After a node generates the `TRAFFIC_SUMMARY`, it should reset the counters that it maintains. This will allow testing of the software for multiple runs.

Upon receipt of the `TRAFFIC_SUMMARY` from all the registered nodes, the registry will proceed to print out the table as depicted below. Each row must be printed on a separate line.

Example output at the registry:

The collated outputs from 10 nodes are depicted below. Note how the number of received messages may be slightly different than the number of sent messages. The summation of sent or received messages at a node may be negative. In this particular example the final summation across all nodes is positive, it may well be negative in your case and that is fine!

	Number of messages sent	Number of messages received	Summation of sent messages	Summation of received messages	Number of messages relayed
Node 1	25000	25440	-340,040,800,604.00	-144,703,367,090.00	40445
Node 2	25000	25395	277,777,554,744.00	192,844,494,434.00	55435
Node 3	25000	24535	-42,851,633,614.00	199,699,309,204.00	60770
Node 4	25000	25130	184,871,797,810.00	91,406,191,639.00	30535
Node 5	25000	24245	-106,636,042,422.00	-180,588,270,287.00	10140
Node 6	25000	25120	24,251,523,172.00	398,033,468,762.00	78545
Node 7	25000	25205	145,053,292,085.00	-377,484,205,221.00	45675
Node 8	25000	24280	-235,398,166,411.00	51,922,993,583.00	8765
Node 9	25000	24985	-70,572,398,997.00	-100,564,359,421.00	15655
Node 10	25000	25665	328,837,533,087.00	34,726,403,247.00	16560
Sum	250000	250000	165,292,658,850.00	165,292,658,850.00	

This is very important to ensure effective grading: The registry must print the table using spaces to separate the columns and new lines to separate the rows. The strings Node 1, Node 2, and so one must be replaced with the corresponding <ip:port> values. The header listed above is for descriptive purposed and is not required. Here is an example:

```
192.168.1.10:5001 25000 25440 -340040800604.00 -144703367090.00 40445
192.168.1.11:5002 25000 25395 277777554744.00 192844494434.00 55435
192.168.1.12:5003 25000 24535 -42851633614.00 199699309204.00 60770
192.168.1.13:5004 25000 25130 184871797810.00 91406191639.00 30535
192.168.1.14:5005 25000 24245 -106636042422.00 -180588270287.00 10140
192.168.1.15:5006 25000 25120 24251523172.00 398033468762.00 78545
192.168.1.16:5007 25000 25205 145053292085.00 -377484205221.00 45675
192.168.1.17:5008 25000 24280 -235398166411.00 51922993583.00 8765
192.168.1.18:5009 25000 24985 -70572398997.00 -100564359421.00 15655
192.168.1.19:5010 25000 25665 328837533087.00 34726403247.00 16560
sum 250000 250000 165292658850.00 165292658850.00
```

5 Command line arguments for the two components

Your classes should be organized in a package called `csx55.overlay`. The command-line arguments and the order in which they should be specified for the Messaging node and the Registry are listed below

```
java csx55.overlay.node.Registry portnum
```

```
java csx55.overlay.node.MessagingNode registry-host registry-port
```

6 Grading

Homework 1 accounts for 10 points towards your final course grade.

Every command that you successfully and correctly implement is worth the following points listed below:

list-messaging-nodes	1pt
list-weights	1 pt
setup-overlay <number-of-connections>	1 pt
send-overlay-link-weights	1 pt
start <number-of-rounds>	3 pt
print-mst	2 pt
exit-overlay	1 pt

Automatic grading of assignments

We have designed a framework for automated grading of assignments in a distributed environment. Once you submit your assignment solution, the framework: (1) crawls Canvas, (2) retrieves your submission, (3) subjects your assignment to a suite of tests, and (4) puts your score alongside comments into Canvas. In most cases, you will see your scores on the programming assignments in about 45-60 seconds.

This will allow you to hone your submissions to ensure a full-score on the programming assignments. You have unlimited attempts (till the submission deadline), and your highest score will be retained. The key idea is that it is now possible for you to know your score before the assignment submission deadline. Use of these autograders is predicated on you following the outputs as specified in the assignment. This requirement isn't arbitrary. In the real world, too, programs are often daisy-chained together, and precise outputs matter. If you are struggling to match the prescribed outputs format, please get in touch with the TAs.

Don't procrastinate ... start early and give yourself room to fix small issues before they become big ones.

7 Deductions

There will be a **10-point deduction** if any of the restrictions below are violated.

1. The data that you will be sending will be `byte[]`. None of your classes can implement the `java.io.Serializable` interface.
2. No GUIs should be built under any circumstances. These are auxiliary paths and the deduction is in place to ensure that none of you attempt to do this.

8 Milestones:

You have 4 weeks to complete this assignment. The weekly milestones below correspond to what you should be able to complete at the end of every week. Please do not wait until the last week to submit. Even the best students have to submit about 4 times to get it right. Feel free to submit partial solutions and improve over the feedback you receive.

Milestone 1: You should be able to have two nodes talking to each other i.e. you are able to exchange messages between two servers.

Milestone 2: You should be able to have 10 messaging node instances talk to the registry, and have the registry sending commands to orchestrate the setting up of the overlay and link weights. You should also be able to issue all commands at the foreground processes.

You should have submitted a version of your assignment at least once by this point; this allows you to see how the autograders have graded your assignment.

Milestone 3: You should be able to compute the MST at each messaging node, generate and enforce routing plans to route messages fed into the overlay. You should be able to track the summation counts for the messages and the contents of these messages.

You should have submitted a version of your assignment at least 3 times by this point.

Milestone 4: Iron out any wrinkles that may preclude you from getting the correct (i.e. not corrupted) outputs at all times.

9 What to Submit

Use **CANVAS** to submit a single .tar file that contains:

- all the Java files related to the assignment (please document your code)
- the `build.gradle` file you use to build your assignment
- a `README.txt` file containing a manifest of your files and any information you feel the TAs needs to grade your program.

Do not include any additional files or folders.

Software versioning: Java 11 and gradle version 8.3

This environment is provided on CS lab machines using module load in Bash:

```
module load courses/cs455
```

```
module load courses/cs555
```

Filename Convention: The class names for your messaging node and registry should be as specified in Section 5. You may call your support classes anything you like. All classes should reside in a package called `csx55.overlay`. The archive file should be named as `<FirstName>_<LastName>_HW1.tar`. For example, if you are Cameron Doe then the tar file should be named `Cameron-Doe-HW1.tar`.

10 Version Change History

This section will reflect the change history for the assignment. It will list the version number, the date it was released, and the changes that were made to the preceding version. Changes to the first public release are made to clarify the assignment; the spirit or the crux of the assignment will not change.

Version	Date	Comments
1.0	8/25/2025	First public release of the assignment.