

Homework 3

IMPLEMENTING THE PASTRY PEER TO PEER NETWORK

VERSION 1.1

DUE DATE: Wednesday October 29th, 2025 @ 8:00 pm

OBJECTIVE

The objective of this assignment is to build a simple peer to peer network where individual peers have 16-bit identifiers. This assignment has several sub-items associated with it: this relates to constructing the logical overlay and traversing the network efficiently to store and retrieve content. This assignment will be modified to clarify any questions that arise, but the crux of this assignment will not change.

All communications in this assignment are based on **TCP**. The assignment must be implemented in **Java** and no external jar files or libraries are allowed. You must develop all functionality yourself. This assignment may be modified to clarify any questions (and the version number incremented), but the crux of the assignment and the distribution of points will not change. You are required to work alone on this assignment. Use of GenAI tools is expressly prohibited; see the textbox below. This assignment will account for **10 points** towards your cumulative course grade. There are several components to this assignment, and the points-breakdown is listed in the remainder of the text. This assignment is to be done individually.

Generative AI Use and Consequences

Use of AI tools such as ChatGPT, Claude, Github Co-Pilot, or anything of their kind to write or "improve" your code or written work at *any* stage is prohibited; this includes the ideation phase. It is your responsibility to ensure that you don't have the Github Co-Pilot extension installed in your IDE; assignment solutions generated by Co-Pilot aren't written by you. Turning in code or an essay written by generative AI tools will be treated as turning in work created by someone else, namely an act of plagiarism and/or cheating. At a minimum, this will result in a 100% deduction (i.e., you will receive a -10/10). To ensure fairness and maintain integrity, grading will also include code reviews, interviews, and on-the-spot code modifications.

Ultimately, you will get out of the class what you put in. Simply copying and pasting code from generative AI tools is not only unethical, it robs you of the chance to learn. Here are four reasons why these generative AI tools undercuts your own education:

1. They take away the struggle that leads to understanding. They rob you of the ability to think and learn the concepts for yourself. Solving problems yourself is how concepts stick. If the AI does the work, what's left for you to learn?
2. You will struggle with the in-classroom quizzes and exams where you will not have access to these tools.
3. Yes, AI tools will become an important part of a software engineer's workflow. But to use them effectively later, you first need solid expertise in the subject matter; and, that only comes from practicing *without* them.
4. These tools are prone to generating imperfect or even incorrect solutions, so trusting them blindly can lead to bad consequences.

Peer Identifiers

Each peer has a 16-bit identifier i.e. the total number of peers in the system can be about 64,000. The system should support specification (at the command line) of an identifier during startup of peer. Note the identifiers should be based on hexadecimals; they hexadecimal identifiers harnessed during routing. I have attached my Java code for converting a Hex String into a byte[] and vice-versa in the appendix portion of this assignment. Use this also for printing out the entries in the routing table and so forth. This feature of assigning identifiers statically to nodes will be used during the scoring process.

1 The Discovery Node

There will also be a discovery node in the system that maintains information about the list of peers in the system. Every time a peer joins or exits the system it notifies this discovery node. The registration information includes information about the peer such as:

- Its 16-bit identifier
- The $\{host:port\}$ information (please use TCP for communications)
- A nickname for this node

In the unlikely case that there is a collision in the ID space, the discovery node notifies the peer about this collision at which point the peer will regenerate a new identifier and repeat the process.

The discovery node has been introduced here to simplify the process of discovering the first peer that will be the entry point into the system. The discovery node is ONLY responsible for

1. Returning **ONE** random node from the set of registered nodes
2. Detect collisions

If the discovery node is used for anything else there will be a **9 point deduction**. An example of misusing the discovery node is to use it to give a new node information about all nodes in the system: such a misuse will defeat the purpose of this assignment.

2 Protocol for Routing Content in the P2P Network

The primary functionality provided by the DHT is the lookup operation. A *lookup(key)* operation identifies the node with the numerically closest identifier to the key. The routing algorithm for the DHT involves two data structures that assist in routing: (1) Leaf Set and (2) Routing table.

Routing can be done using just the Leaf Set; though the routing solution converges in this case, the solution is inefficient. The best routing solution combines both the Leaf Set and the Routing Table.

Your implementation should support the solution that combines the Leaf Set and the Routing Table. The description that follows in the remainder of the assignment is based on the Pastry algorithm as described in our recommended text.

2.1 Leaf Set

At a given peer, this data structure is responsible for tracking the $2l$ **neighbors** of that peer; l to the right of the peer and l to its left. The DHT ID space can be thought of as being organized as a ring: $0-(2^{16}-1)$. The neighbors refer to peers whose identifiers are *numerically closest* to the peer in question.

So, if there is a set of peers in the system with IDs: 53, 65, 69, 73, 83, 92. Then the Leaf Set at 69, when $l=2$, is {53, 65, 73, and 83}.

The simplest routing solution using the Leaf Set involves taking hops (of size >0 and $\leq l$) to reach the destination. At each peer, you will choose a hop that gets you close to the destination; so, in most cases, you will be taking l hops at each peer as you try to get closer to the destination, before using a hop of size 1. For a system with N peers, it will require about $N/2^l$ hops to deliver a message using only the Leaf Set (as can be seen, this is very inefficient).

For the purposes of this assignment $l=1$ i.e. your Leaf Set at each peer involves 2 (numerically) closest peers one with ID greater and one with ID lower than the peer.

2.2 The Routing Table

The routing table maintains information about several peers in the system. All peer identifiers in the DHT are viewed as **hexadecimal** values. The routing table *classifies* peer identifiers based on their hexadecimal **prefixes**. The table has *as many rows* as there are hexadecimal digits in the peer identifier. Thus, in our case, with 16-bit identifiers, there will be $16/4 = 4$ rows. Any row in this table contains 16 entries: 1 for each possible value of the n^{th} hexadecimal digit. The routing table *excludes entries/values* that correspond to local peer's identifier. As can be seen the routing table (depicted in Figure 1) at a peer, has increased density of coverage for peers with IDs that are numerically closer to its own. The "n" in the cells refers to node handles; each cell contains the IP address and identifier of *one* peer that matches the prefix criteria. Note that some of the cells may be empty since the Routing Table may not be fully populated.

p =	Identifier Prefixes and corresponding node handles n															
0	0 n	1 n	2 n	3 n	4 n	5 n	6 n	7 n	8 n	9 n	A n	B n	C n	D n	E n	F n
1	60 n	61 n	62 n	63 n	64 n	65 n	66 n	67 n	68 n	69 n	6A n	6B n	6C n	6D n	6E n	6F n
2	65 0 n	65 1 n	65 2 n	65 3 n	65 4 n	65 5 n	65 6 n	65 7 n	65 8 n	65 9 n	65 A n	65 B n	65 C n	65 D n	65 E n	65 F n
3	65 A0 n	65 A1 n	65 A2 n	65 A3 n	65 A4 n	65 A5 n	65 A6 n	65 A7 n	65 A8 n	65 A9 n	65 AA n	65 AB n	65 AC n	65 AD n	65 AE n	65 AF n

Figure 1: Example routing table at a node with identifier **65A1**. Each cell contains the IP address and identifier of one peer that matches the prefix criteria. Note that some of the cells may be empty.

The routing solution that uses the Routing Table and the Leaf Set provides the best routing characteristics. Specifically, the algorithm routes requests to any node in $O(\log N)$ messages.

Let's say that you receive a message from a peer **A** with the intended destination of **D**. You will begin by comparing the hexadecimal representations for peers **A** and **D** from *left-to-right* to discover, p , their longest common prefix. You will then consult the routing table at **A** using the p (that we just computed) as the row offset. The first non-matching digit in **D** is used as the column offset, to access the required element in the table. The table construction should ensure that this cell contains the address of the peer that has $(p + 1)$ prefixes in common with **D**.

Another check that is performed is if **D** is in the leaf set of **A**. If it is, the routing converges and after routing to the destination **D** no further steps need to be performed.

2.3 Addition of a New Peer and Construction of the Routing Table

New nodes use a joining protocol. This protocol allows a new peer to *acquire* its routing table and leaf set contents; this protocol also includes notifying other nodes of changes that they must make to their tables.

Let's say that a new peer joins the system and this new node's identifier is **X**. This new node contacts its entry-point node **A** using the Discovery Node. Node **X** sends a special join request message to **A** and gives **X** as its destination. This node **A** dispatches the join message via the DHT to an existing node with an identifier that is numerically closest to **X**; we will call this the destination node **Z**.

The join message is routed through the network: **A**, **Z** and intermediate nodes (**B**, **C**, ...). This results in the *transmission of relevant parts* of their routing tables and leaf sets to **X**. **X** examines and constructs its own routing table and leaf set from the: entry point peer, intermediate peers, and the destination peer.

In the original protocol, the node that serves as the entry (**A**) has network proximity to the new node **X**. We will be running our assignment in a local cluster, so this technically moot in our case. However, similar to the original algorithm, we will assume that the first row of **A**'s table is a good initial choice for **X** i.e. **A** and **X** will have the same first row. Note: the Routing Tables uses 0-based row indexing i.e. the indices start at 0.

A's table is not relevant for the second row because the GUIDs for **X** and **A** may not share the 1st digit. But the routing algorithm ensures that **X** and **B**'s GUID do share the first hexadecimal digit. So the second row of **B**'s routing table B_1 is a suitable initial value for X_1 . Similarly, C_2 is suitable for X_2 and so on.

Since **Z**'s identifier is numerically closest to **X**'s. **X**'s ideal leaf set will differ from **Z**'s by *just one* member. This is **eventually** optimized through interaction with the neighbors.

Once **X** has constructed its leaf set and routing table, **X** sends its contents to all *nodes* identified in the **leaf set** and the **routing table**. The nodes that receive these updates, *adjust* their own tables to **incorporate** the node.

3 Storing data items

You must use complete routing solution (encompassing the routing table and the leaf set) to store data items at the appropriate node. A data item with a key k will be stored at the peer with the *closest* numerical identifier (in the case of ties choose the peer with the higher identifier). The data item that will be given to you will include images, text, and other types. To support this feature, you will develop a StoreData program that accepts as input the file that needs to be stored. This StoreData program contacts a random peer (you can contact the Discovery node to retrieve this information). The StoreData program will first compute the 16-bit digest (appendix A) for the file name and then use this hash to **lookup** the peer where it should be stored: you will be contacting the aforementioned random peer to initiate this lookup; the node that gets back to you will be the node that is most suitable to store your data. The file is then transferred to that suitable peer, which is responsible for storing the file in the `/tmp/<peer-id>` directory of the machine that it is running on.

4 Diagnostics

To make sure that things are progressing correctly this assignment requires that several diagnostic information be printed on the console. These include:

1. The routing table and the leaf set at a node
2. The list of files managed by the node: This would be stored in the `/tmp` directory of the machine on which the peer is running
3. Print out a message every time you route a query: This message should indicate the hop number that it corresponds to in the routing path. So, if a `lookup()` operation has bounced off of 3 nodes, and is now received at a node ... it should print a hop count of 4.
 - a. We will use this information to reconstruct the path which the lookup/successor operations took.
 - b. In the absence of this hop information, all that would need to be done is to sort the peer ids appropriately to simulate the correct path.

5 Third-party libraries and restrictions:

You are not allowed to download *any* other code from *anywhere* on the Internet. You are also not allowed to use RPC or distributed object frameworks to develop this functionality (there is a **10 point deduction** for this). You should not build GUIs for this application; in the context of this assignment GUI-building is an auxiliary path (there is a **10 point deduction** for building a GUI). You can discuss the project with your peers at the architectural level, but the project implementation is an individual effort.

6 Testing Scenario

Your submission will be tested with between 10-20 nodes possibly on different machines. The port number on which your peer runs and listens to for communications should be configurable.

7 Programs, Commands, and Points

You must implement the following three programs with the corresponding commands. Follow the given formats. In the commands, all strings formatted as `<example>` are replaced with the corresponding value. For example, `<port>` can be replaced with the value 12345.

Discover Node

This program keeps track of the peers in the ring and returns a random peer to new nodes.

Start:

```
java csx55.pastry.node.Discover <port>
```

Commands and example outputs:

list-nodes

```
12.34.56.78:4821, a3f2  
201.122.8.199:17345, 4c9b  
77.88.99.11:8020, 1d7e  
153.6.240.4:6550, 9b01
```

Note that the command prints ips, ports, and ids of every node currently in the ring. **1pt**

Peer Node

This program is a peer in the ring.

Start:

```
java csx55.pastry.node.Peer <discover-host> <discover-port> <id>
```

The id is specified in hexadecimal

Commands and example outputs:

id

```
47ab
```

This prints the id of the node. **1pt**

leaf-set

```
23.45.67.89:3344, 1a2b  
198.51.100.23:5500, 4f3c  
65.77.88.99:9201, 9d0e  
172.16.5.4:8088, be7f
```

This prints the leaf-set for the node. This includes ip, port, and id for each entry. The leaf-set is sorted by id. **1pt**

routing-table

```
0-:,1-110.50.67.25:57414,2-:,3-130.207.99.82:60275,4-62.182.9.150:12939,5-:,6-97.104.137.136:27330,7-  
155.146.28.146:7495,8-238.201.43.107:31197,9-30.122.170.145:58399,a-79.117.55.241:12060,b-:,c-  
164.35.113.175:16229,d-244.177.7.40:19705,e-245.109.105.22:39855,f-41.117.89.52:45487  
60-:,61-:,62-:,63-:,64-:,65-:,66-:,67-:,68-:,69-:,6a-:,6b-:,6c-:,6d-97.104.137.136:27330,6e-:,6f-:
```

```
6d0-97.104.137.136:27330,6d1-:,6d2-:,6d3-:,6d4-:,6d5-:,6d6-:,6d7-:,6d8-:,6d9-:,6da-:,6db-:,6dc-:,6dd-:,6de-:,6df-:
6d00-:,6d01-:,6d02-:,6d03-:,6d04-97.104.137.136:27330,6d05-:,6d06-:,6d07-:,6d08-:,6d09-:,6d0a-:,6d0b-:,6d0c-:,6d0d-:,6d0e-:,6d0f-:
```

This prints the full routing table for the node. Follow section 2.2. The routing table is comprised of 4 rows. Rows are separated by the new line character. Each row contains 16 cells, cells are comma-separated, each cell contains <id>-<ip>:<port>. Leave ip and port empty if there are no values. **1pt**

list-files

```
test.png, 9c2b
file.txt, 9c65
```

This lists all files currently stored on this node. **1pt**

exit

This command triggers an exit from the ring and kills the peer node. **1pt**

DataStorage Program

The DataStorage program executes one store/retrieve operation and then terminates immediately.

Start command:

```
java csx55.pastry.node.Data <discover-host> <discover-port> <mode> <path>
```

There are two possible modes, store and retrieve.

Example of storing a file:

```
java csx55.pastry.node.Data denver 12345 store ./images/test.png
```

The image is located at ./images/test.png, but only the name test.png is used for the item key to store the image. The image must be stored by the correct peer node. The peer node must store it inside a folder <peer-id>/.

The output of this program must be the list of ids of all nodes on the route followed by the id of the item being stored.

For example:

```
1a3f
2b4c
3e1d
7f0a
9c2b
```

In this case the route is 1a3f -> 2b4c -> 3e1d -> 7f0a and the id of the image is 9c2b. **2pt**

Example of downloading a file:

```
java csx55.pastry.node.Data denver 12345 retrieve ./downloads/test.png
```

The images must be downloaded to ./downloads/test.png, but only the name test.png is used for the item key to retrieve the image.

The output of this program is the same as the output for storing a file. **2pt**

Restrictions

- Programs should not be launched from an IDE such as Eclipse.
- Java object serialization should not be used when storing files in the file system.

General Recommendations

- It is recommended to use `java.util.logging` for printing logs instead of using standard system output.
- It is encouraged to create scripts to automate tasks such as compiling source code and launching programs.
- When you are launching processes on 25 machines, please do not login to each one manually ... use a script instead.

8 What to Submit

Use **CANVAS** to submit a single .tar file that contains:

- all the Java files related to the assignment (please document your code)
- the `build.gradle` file you use to build your assignment
- a `README.txt` file containing a manifest of your files and any information you feel the TAs needs to grade your program.

Software versioning: Java 11 and Gradle version 8.3

This environment is provided on CS lab machines using module load in Bash:

```
module load courses/cs455
```

```
module load courses/cs555
```

Filename Convention: The class names for your peer node and discover nodes should be as specified in Section 7. You may call your support classes anything you like. All classes should reside in a package called `csx55.pastry`. The archive file should be named as `<FirstName>_<LastName>_HW3.tar`. For example, if you are Cameron Doe then the tar file should be named `Cameron-Doe-HW3.tar`.

9 Version Change History

Version	Date	Change
1.1	10/2/2025	Example outputs have been expanded.
1.0	10/1/2025	First release of the assignment with Appendix A and Appendix B.

Appendix A

```
/**
 * This method converts a set of bytes into a Hexadecimal representation.
 */
public String convertBytesToHex(byte[] buf) {
    StringBuffer strBuf = new StringBuffer();

    for (int i = 0; i < buf.length; i++) {
        int byteValue = (int) buf[i] & 0xff;
        if (byteValue <= 15) {
            strBuf.append("0");
        }
        strBuf.append(Integer.toString(byteValue, 16));
    }
    return strBuf.toString();
}

/**
 * This method converts a specified hexadecimal String into a set of bytes.
 */
public byte[] convertHexToBytes(String hexString) {
    int size = hexString.length();
    byte[] buf = new byte[size / 2];

    int j = 0;
    for (int i = 0; i < size; i++) {
        String a = hexString.substring(i, i + 2);
        int valA = Integer.parseInt(a, 16);

        i++;

        buf[j] = (byte) valA;
        j++;
    }

    return buf;
}

/**
 * This method computes the 16-bit digest of a file name.
 */
public static byte[] hash16(String fileName) {
    int h = fileName.hashCode() & 0xFFFF;
    byte[] result = new byte[2];
    result[0] = (byte) ((h >>> 8) & 0xFF);
    result[1] = (byte) (h & 0xFF);
    return result;
}
```